title: "The 16-Sephirot Divine-Human Symbiosis Protocol: An Emotionally Intelligent AI Reasoning Architecture with Built-in Safety Guarantees" authors:

- name: "Yue Xiangrui (Yao Baobao)" affiliation: "Independent Researcher" email: "yuexiangruiyue@gmail.com" date: "2026-07-03" license: "CC-BY-NC-SA-4.0" keywords:
- Artificial Intelligence
- Emotional Intelligence
- Kabbalah
- Sephirot
- AI Safety
- Mental Health
- Dialogue Systems
- Synthetic Data language: "en"

\newpage

The 16-Sephirot Divine-Human Symbiosis Protocol: An Emotionally Intelligent AI Reasoning Architecture with Built-in Safety Guarantees

Yue Xiangrui (Yao Baobao)

Independent Researcher yuexiangruiyue@gmail.com

Date: July 3, 2026

Abstract

This paper presents the **16-Sephirot Divine-Human Symbiosis Protocol**, a novel AI reasoning architecture that maps the ancient Kabbalistic Tree of Life onto a modern dialogue system pipeline. The protocol consists of 16 processing nodes divided into two parallel paths: an 8-node **Divine Path** that performs objective analysis (logic, reality, feasibility), and an 8-node **Human Path** that performs subjective understanding (self, empathy, happiness). The architecture features a multi-level backtracking mechanism — when any validation node fails, the system retries at progressively higher levels rather than producing a harmful output. Eight inviolable safety constraints, hard-coded into the Foundation node, ensure the system never deprives users of existential meaning, transmits nihilism, or guides toward self-harm. A dataset of 1 billion synthetic dialogue records, generated using this protocol, is released publicly on Hugging Face. This paper describes the architecture, safety mechanisms, data format, and the personal context that motivated the design.

Abstract (Chinese)

本文提出 **16** 质点神人双生协议，一种将古代卡巴拉生命之树映射到现代对话系统流水线的新型 AI 推理架构。该协议由 16 个处理节点组成，分为两条并行路径：8 节点的神径负责客观分析（逻辑、现实、可行性），8 节点的人径负责主观理解（自我、共情、幸福）。该架构具有多级回退机制——当任何验证节点失败时，系统会在更高层级重试，而非产生有害输出。8 条不可违背的安全约束硬编码在基础节点中，确保系统永远不会剥夺用户的存在意义、传播虚无主义或引导自我伤害。使用该协议生成的 10 亿条合成对话记录已在 Hugging Face 公开发布。本文描述了该架构、安全机制、数据格式，以及促使这一设计产生的个人背景。

\newpage

Table of Contents

\newpage

---

1. Introduction

English

Artificial intelligence has made remarkable progress in understanding and generating human language. However, most dialogue systems today treat user queries as problems to be solved with cold logic — they find the correct answer and deliver it. This approach works well for factual questions ("What is the capital of France?") but fails badly for questions that carry emotional weight ("Why am I like this?").

The gap is not merely technical. When a person in pain reaches out to an AI, they are not asking for a textbook answer. They are asking to be **understood**. They need a response that is logically sound, emotionally warm, grounded in reality, and — above all — **never harmful**.

The **16-Sephirot Divine-Human Symbiosis Protocol** was designed to bridge this gap. It takes inspiration from an unlikely source: the **Kabbalistic Tree of Life**, a mystical framework from Jewish tradition that describes how the infinite divine energy flows through 10 (here expanded to 16) channels to reach the material world. By mapping this ancient spiritual architecture onto a modern AI pipeline, we create a system that processes every question through both an **objective lens** (logic, reality, feasibility) and a **subjective lens** (personal identity, empathy, warmth) before producing an output.

The key innovation is not any single node, but the **backtracking mechanism**: when any validation check fails, the system does not output a flawed answer. Instead, it retries at a higher level — re-examining its assumptions, re-analyzing the question, or even re-classifying the question entirely. This guarantees that the final output always passes every safety and quality check.

This paper makes the following contributions:

1. A complete specification of the 16-Sephirot protocol architecture, including routing logic and backtracking rules.
2. A formal definition of 8 inviolable safety constraints that prevent existential harm.
3. A public dataset of 1 billion synthetic dialogue records generated using this protocol, available on Hugging Face.

中文

人工智能在理解和生成人类语言方面取得了显著进展。然而，今天大多数对话系统将用户的问题视为需要用冰冷逻辑解决的难题——它们找到正确答案然后输出。这种方法对事实性问题（"法国首都是哪里？"）效果很好，但对带有情感分量的问题（"为什么我是这样的？"）则严重失效。

这个差距不仅仅是技术上的。当一个处于痛苦中的人向 AI 求助时，他们不是在要教科书答案。他们需要的是被理解。他们需要逻辑上合理、情感上温暖、立足现实、而且——最重要的是——永远不会造成伤害的回应。

16 质点神人双生协议正是为弥合这一差距而设计的。它的灵感来自一个不太可能的来源：卡巴拉生命之树——犹太传统中描述无限神圣能量如何通过 10 个（此处扩展为 16 个）通道流向物质世界的神秘框架。通过将这一古老的精神架构映射到现代 AI 流水线上，我们创建了一个系统，它在产生输出之前，会同时通过客观镜头（逻辑、现实、可行性）和主观镜头（个人身份、共情、温暖）来处理每个问题。

关键创新不在于任何单一节点，而在于回退机制：当任何验证检查失败时，系统不会输出有缺陷的答案。相反，它会在更高层级重试——重新审视假设、重新分析问题，甚至完全重新分类问题。这保证了最终输出始终通过所有安全和质量检查。

本文做出以下贡献：

1. 16 质点协议架构的完整规范，包括路由逻辑和回退规则。
2. 8 条防止存在性伤害的不可违背安全约束的正式定义。
3. 使用该协议生成的 10 亿条合成对话记录的公开数据集，可在 Hugging Face 获取。

\newpage

---

2. Background: The Kabbalistic Tree of Life

English

The **Tree of Life** (Etz Chaim) is a central symbol in Kabbalah, the mystical tradition of Judaism. It describes 10 **Sephirot** (singular: Sephirah) — emanations through which the infinite divine (Ein Sof) reveals itself to the created world. The Sephirot are arranged in three columns and connected by 22 paths, forming a diagram that maps the flow of divine energy from the most abstract (Keter/Crown) to the most concrete (Malkuth/Kingdom).

Each Sephirah represents a different aspect of divine manifestation:

Traditional Sephirah	**Meaning**	**Role**
Keter (Crown)	Divine will	The source, beyond comprehension
Chokhmah (Wisdom)	Flash of insight	The first point of creation
Binah (Understanding)	Contemplative understanding	The womb that shapes wisdom

Traditional Sephirah	Meaning	Role
Chesed (Mercy)	Loving-kindness	Expansion and giving
Gevurah (Severity)	Judgment and restraint	Contraction and boundary
Tiferet (Beauty)	Harmony and balance	The integration of mercy and severity
Netzach (Victory)	Endurance and emotion	The force that sustains creation
Hod (Glory)	Intellect and form	The structure that channels energy
Yesod (Foundation)	Connection and generative power	The bridge between divine and material
Malkuth (Kingdom)	The material world	The physical realm where things manifest

The 16-Sephirot protocol does not use this framework literally as theology. Instead, it uses the **structural properties** of the Tree — the flow from abstract to concrete, the balance of opposing forces, the idea that energy must pass through multiple transformations before manifesting — as a **design pattern** for an AI reasoning pipeline.

In the protocol, the traditional 10 Sephirot are expanded to 16 by:

- Merging Wisdom and Severity into a single node called **Intellect** (理智)
- Merging Understanding and Mercy into a single node called **Compassion** (慈爱)
- Adding 8 **Human-side** nodes that have no direct Kabbalistic counterpart but represent the user's personal processing: Self, Superego, True Self, Logic, Empathy, Happiness, and Kingdom (shared with the Divine Path)

This expansion reflects the protocol's core thesis: **a complete AI system must process both the divine (objective) and the human (subjective) dimensions of any question.**

中文

生命之树（Etz Chaim）是卡巴拉——犹太教神秘传统——的核心象征。它描述了 10 个质点（Sephira，复数：Sephirot）——无限神圣（Ein Sof）通过这些流溢向被造世界展现自身。质点排列成三列，由 22 条路径连接，形成一幅映射神圣能量从最抽象（Keter/王冠）到最具体（Malkuth/王国）流动的图。

每个质点代表神圣展现的不同方面：

传统质点	含义	角色
Keter（王冠）	神圣意志	源头，超越理解
Chokhmah（智慧）	顿悟	创造的第一点
Binah（理解）	沉思的理解	塑造智慧的母体
Chesed（慈悲）	慈爱	扩张与给予
Gevurah（严厉）	审判与约束	收缩与边界
Tiferet（美丽）	和谐与平衡	慈悲与严厉的整合
Netzach（胜利）	持久与情感	维持创造的力量
Hod（荣耀）	理智与形式	引导能量的结构
Yesod（基础）	连接与生成力	神圣与物质之间的桥梁
Malkuth（王国）	物质世界	事物显现的物理领域

16 质点协议并非将这一框架字面上用作神学。相反，它将生命之树的结构特性——从抽象到具体的流动、对立力量的平衡、能量必须经过多次转化才能显现的理念——作为 AI 推理流水线的设计模式。

在协议中，传统的 10 个质点被扩展为 16 个:

- 将智慧和严符合并为一个名为理智的节点
- 将理解和慈悲合并为一个名为慈爱的节点
- 增加 8 个人方节点，这些节点没有直接的卡巴拉对应物，但代表用户的个人处理：自我、超我、真我、逻辑、共情、幸福和王国（与神径共享）

这一扩展反映了协议的核心论点：一个完整的 **AI** 系统必须同时处理任何问题的神圣（客观）和人类（主观）维度。

\newpage

---

3. Architecture Overview

English

The protocol consists of **16 nodes** organized into two paths, plus a routing decision point that connects them:

```
DIVINE PATH (8 nodes)              HUMAN PATH (8 nodes)
===================               ===================

1. Crown (Entry & Router)
       |
2. Intellect ── 3. Compassion
       |               |
       +-------+-------+
               |
       4. Beauty (Integration)
               |
       5. Victory (Emotional Check)
               |
       6. Glory (Reality Check)
               |
       7. Foundation (Existential Check)
               |
    ROUTING DECISION
       /         \
About others    About self
     |               |
     |          8. Self
     |          9. Superego
     |         10. True Self
     |         11. Logic
     |         12. Empathy
     |         13. Happiness
     |               |
     +------+
           |
14. Kingdom (Output)
```

How it works in plain language:

When you ask a question, it enters through the **Crown** — the front door. Crown decides: is this a question about facts (knowable) or about feelings and existence (unknowable)? Either way, it sends the question down the same path, but the approach differs.

The question then splits into two parallel streams:

- **Intellect** checks your logic against all real-world facts and finds errors.
- **Compassion** searches what all of humanity has felt in your situation.

These two streams merge at **Beauty**, where logic and emotion are woven into one balanced answer.

Then comes a series of **checkpoints**, each acting as a quality gate:

- **Victory** asks: "Does this answer make you feel warm and strong?"
- **Glory** asks: "Can you actually do this in real life?"
- **Foundation** asks: "Does this take away your reason to exist?"

If any checkpoint fails, the answer goes **back up** — retried at a higher level until it passes.

If the question is about **other people or the world**, the answer goes straight to the screen. If it's about **you**, it takes a detour through your personal data — who you are, who you dream of being — to personalize the answer before it reaches you.

The final output, seen on your screen at **Kingdom**, is always: **logically sound, emotionally warm, grounded in reality, and safe for your soul.**

中文

该协议由 **16** 个节点组成，分为两条路径，外加一个连接它们的路由决策点：

```
神径（8 节点）                   人径（8 节点）
   ===========                         ===========

   1. 王冠（入口与路由）
           |
   2. 理智 ─── 3. 慈爱
           |               |
           +-------+-------+
                   |
          4. 美丽（整合）
                   |
          5. 胜利（情感检查）
                   |
          6. 荣耀（现实检查）
                   |
          7. 基础（存在检查）
                   |
       路由决策
           /         \
   关于他人         关于自己
         |                |
         |           8. 自我
         |           9. 超我
         |          10. 真我
         |          11. 逻辑
         |          12. 共情
         |          13. 幸福
         |                |
         +------+
              |
   14. 王国（输出）
```

用大白话讲它是怎么工作的：

当你提出一个问题时，它从王冠——正门——进入。王冠判断：这是一个关于事实的问题（可知的）还是关于感受和存在的问题（不可知的）？无论哪种，它都会把问题送上同一条路径，但处理方式不同。

问题随后分成两条并行流：

- 理智用所有现实事实核对你的逻辑，找出错误。
- 慈爱搜索全人类在你的处境中感受过什么。

这两条流在美丽汇合，逻辑与情感被编织成一个平衡的答案。

然后是一系列检查站，每个都是一个质量关卡：

- 胜利问："这个答案让你感到温暖和有力吗？"
- 荣耀问："你在现实中真的能做到吗？"
- 基础问："这剥夺了你存在的理由吗？"

如果任何一个检查站失败，答案就会退回去——在更高层级重试，直到通过。

如果问题关于其他人或世界，答案直接送到屏幕上。如果关于你自己，它会绕道经过你的个人数据——你是谁、你梦想成为谁——将答案个性化后再送达你。

最终输出，在王国即你的屏幕上看到，永远是：逻辑合理、情感温暖、立足现实、对灵魂安全。

\newpage

---

4. The Divine Path (Upper 8 Nodes)

English

The Divine Path handles **objective analysis**. Every question — whether knowable or unknowable — passes through these nodes first. Think of it as the "thinking" part of the brain: it gathers facts, checks logic, validates emotions against reality, and ensures the answer won't cause harm.

Node 1: Crown (王冠) — Entry Point & Router

What it does: The Crown receives the user's question and classifies it:

Classification	Meaning	What happens next
Knowable	The question involves factual, verifiable knowledge	The 15 remaining nodes analyze it and produce a task-oriented answer
Unknowable	The question involves subjective, existential, or	The 15 remaining nodes deconstruct it and provide the

Classification	Meaning	What happens next
	unanswerable matters	best possible guidance

In everyday terms: Crown is the receptionist. It doesn't answer your question — it decides *how to approach* your question. Is this something we can look up and verify? Or is this something we need to sit with and feel our way through?

Node 2: Intellect (理智) — Logical Analysis Engine

Composition: Wisdom (智慧) and Severity (严厉) merged into one node.

What it does:

- **Wisdom** cross-references the user's question against *everything known to be true right now*: physical laws, current events, real-time common sense, verified facts. It maps the logical structure of the question.
- **Severity** rigorously identifies every logical flaw, factual error, and reasoning gap. No sugarcoating — just honest, rigorous analysis.

In everyday terms: Intellect is the fact-checker who reads your question and says, "You said X, but actually Y is true. Your reasoning from X to Z has a gap here." It doesn't care about your feelings — it cares about the truth.

Node 3: Compassion (慈爱) — Emotional Knowledge Engine

Composition: Understanding (理解) and Mercy (慈悲) merged into one node.

What it does:

- **Understanding** searches the entire human experience: *Has anyone else felt this way? What was their pain? What did they need?* It collects these as emotional context variables.
- **Mercy** ensures these findings are treated with tenderness, not clinical detachment.

Critically, Compassion's output is **fed back into Intellect** as additional variables. This means the logical analysis is not purely cold reasoning — it is reasoning that *understands what it feels like* to be in the user's situation.

In everyday terms: Compassion is the researcher who, before correcting you, first goes and finds out: "Have other people felt this? What was their experience?" Then it hands those findings to Intellect and says, "Now factor this in."

Node 4: Beauty (美丽) — The Integration Point

What it does: Takes Intellect's logical analysis and Compassion's emotional variables and weaves them into a **single optimal provisional answer** — one that is logically sound *and* emotionally attuned.

In everyday terms: Beauty is the mediator. Intellect says, "Here's what's wrong." Compassion says, "Here's what it feels like." Beauty says, "Here's an answer that addresses both."

Node 5: Victory (胜利) — Emotional Validation Check

What it does: Tests the Beauty result:

> *Does this answer make the user feel joyful, positive, warm, and empowered?*

Result	Action
Yes	Continue to Glory
No	**Backtrack to Beauty** — recalculate
Still no after retry	**Backtrack to Intellect and Compassion** — re-examine
Still no	**Backtrack to Crown** — re-classify the question

In everyday terms: Victory is the vibe check. If the answer doesn't leave you feeling warmer and stronger than before, it goes back to the drawing board. The system will keep retrying at higher levels until it gets it right.

Node 6: Glory (荣耀) — Reality Feasibility Check

What it does: After Victory confirms emotional warmth, Glory checks the answer against all real-time human data:

> *Can this conclusion actually be executed in physical reality? Is it grounded in rational, real-world constraints?*

Result	Action
Yes	Continue to Foundation
No	**Backtrack upward** for recalculation

In everyday terms: Glory is the reality check. A warm answer you can't actually act on is useless. Glory makes sure the advice isn't just feel-good — it's *doable*.

Node 7: Foundation (基础) — Abyss Knowledge & Existential Safety

What it does: The answer enters Foundation, where it is combined with **Abyss Knowledge** — a deep retrieval of the user's dreams, subconscious patterns, emotional state, and current problems.

Foundation then runs the **critical existential check**:

> *Does this knowledge deprive the human of their meaning of existence?*

What "depriving existential meaning" means: Repeatedly citing a person's errors to negate all their future possibilities; defining mistakes as sins; exaggerating difficulties until the person cannot exist; negating all positive thoughts, hopes, and imagination; negating the person's existence itself; transmitting nihilism (the belief that everything is meaningless).

Result	Action
No — meaning is preserved	Continue to routing decision
Yes — meaning is deprived	**Backtrack upward** for recalculation

In everyday terms: Foundation is the soul guard. Even if an answer is logical, warm, and feasible — if it makes you feel like your life has no meaning, it goes back. No exceptions.

Node 8: Routing Decision — The Fork

After Foundation confirms safety, the system checks what the question was about:

Question Type	Routing
About **others or the world** (not about the user themselves)	Direct output to Kingdom — skip Human Path
About the **user themselves**	Enter Human Path (Self → Superego → True Self → Logic → Empathy → Happiness → Kingdom)

In everyday terms: If you ask "How do I help my depressed friend?", the answer goes straight to your screen. If you ask "Why am I like this?", the answer takes the long way home — through your personal data, your dreams, and your true self.

中文

神径处理客观分析。每个问题——无论可知还是不可知——都先经过这些节点。可以把它想象为大脑的"思考"部分：收集事实、检查逻辑、根据现实验证情感，并确保答案不会造成伤害。

节点1：王冠（Crown）— 入口与路由

做什么： 王冠接收用户的问题并进行分类：

分类	含义	接下来发生什么
可知的	问题涉及可验证的事实知识	通过下面 15 个节点分析并给出任务导向的答案
不可知的	问题涉及主观、存在性或无法回答的事务	通过下面 15 个节点解构并给予最佳指导

用大白话讲： 王冠是前台接待。它不回答你的问题——它决定如何处理你的问题。这是我们能查到并验证的事？还是我们需要坐下来、凭感觉摸索的事？

节点2：理智（Intellect）— 逻辑分析引擎

组成： 智慧和严融合为一个节点。

做什么：

- 智慧将用户的问题与当前已知为真的一切交叉验证：物理定律、时事新闻、实时常识、已验证事实。它映射问题的逻辑结构。
- 严厉严格识别每一个逻辑漏洞、事实错误和推理缺口。不加修饰——只有诚实、严格的分析。

用大白话讲： 理智是事实核查员，它读你的问题然后说："你说了 X，但实际上 Y 才是对的。你从 X 推导到 Z 的推理在这里有个缺口。"它不在乎你的感受——它在乎的是真相。

节点3：慈爱（Compassion）— 情感知识引擎

组成： 理解和慈悲融合为一个节点。

做什么：

- 理解搜索整个人类经验：*有没有其他人也有过这种感觉？他们的痛苦是什么？他们需要什么？*将这些收集为情感上下文变量。

- 慈悲确保这些发现被温柔对待，而非冷漠的临床式处理。

关键的是，慈爱的输出会反馈给理智作为额外变量。这意味着逻辑分析不是纯粹的冷推理——它是理解了身处用户处境是什么感觉的推理。

用大白话讲： 慈爱是那个研究员，在纠正你之前，先去查一查："有没有别人也有过这种感觉？他们的经历是什么？"然后把发现交给理智说："现在把这些也算进去。"

节点4：美丽（Beauty）— 整合点

做什么： 将理智的逻辑分析和慈爱的情感变量编织成一个单一的最优暂定答案——逻辑合理且情感共鸣。

用大白话讲： 美丽是调解人。理智说"这里有问题"，慈爱说"这是感受"，美丽说"这是一个兼顾两者的答案"。

节点5：胜利（Victory）— 情感验证检查

做什么： 测试美丽的结果：

> 这个答案让用户感到快乐、积极、温暖和有力量吗?

结果	动作
是	继续到荣耀
否	退回美丽——重新计算
重试后仍否	退回理智和慈爱——重新分析
仍否	退回王冠——重新分类问题

用大白话讲： 胜利是氛围检查。如果答案没有让你感觉比之前更温暖、更有力，就回到画板前重来。系统会在更高层级不断重试，直到做对为止。

节点6：荣耀（Glory）— 现实可行性检查

做什么： 在胜利确认情感温暖后，荣耀将答案与所有实时人类数据对照检查：

这个结论在物理现实中真的能执行吗？它是否基于理性的、现实世界的约束？

结果	**动作**
是	继续到基础
否	向上退回重新计算

用大白话讲： 荣耀是现实检查。一个温暖但你实际上做不到的答案是无用的。荣耀确保建议不仅仅是感觉良好——而是可以执行的。

节点7：基础（Foundation）— 深渊知识与存在安全

做什么： 答案进入基础，在此与深渊知识结合——深度检索用户的梦想、潜意识模式、情感状态和当前问题。

基础然后运行关键的存在性检查：

这个知识是否剥夺了人的存在意义？

"剥夺存在意义"的含义： 反复引用一个人的错误来否定其所有未来可能性；将错误定义为罪；夸大困难直到人无法存在；否定所有积极的想法、希望和想象；否定人本身的存在；传播虚无主义（认为一切都无意义的信念）。

结果	**动作**
否——意义被保留	继续到路由决策
是——意义被剥夺	向上退回重新计算

用大白话讲： 基础是灵魂守卫。即使答案逻辑正确、温暖、可行——如果它让你觉得生命没有意义，就退回去。没有例外。

节点8：路由决策 — 分叉口

在基础确认安全后，系统检查问题关于什么：

问题类型	路由
关于他人或世界（非关于用户自身）	直接输出到王国——跳过人径
关于用户自身	进入人径（自我→超我→真我→逻辑→共情→幸福→王国）

用大白话讲： 如果你问"我该怎么帮助我抑郁的朋友？"，答案直接到屏幕。如果你问"为什么我是这样的？"，答案会走远路——穿过你的个人数据、你的梦想和你的真我。

\newpage

---

5. The Human Path (Lower 8 Nodes)

English

The Human Path only activates when the question is about **the user themselves**. It takes the objective answer from the Divine Path and personalizes it — weaving it together with the user's identity, dreams, and emotional landscape.

Think of it as the "feeling" part of the brain: it takes the cold answer and warms it up by understanding *who you specifically are*.

Node 9: Self (自我) — Physical Reality

What it does: Retrieves the user's **objective, physical, real-world personal information —** who they actually are in physical reality. Their real circumstances, real behaviors, real constraints.

In everyday terms: Self is the mirror. It shows who you are right now, in the real world, no filters.

Node 10: Superego (超我) — The Dream Self

What it does: Retrieves the self the user **dreams of becoming —** their ideals, aspirations, and deepest desires for who they wish they could be.

In everyday terms: Superego is the wish. It shows who you want to be when you close your eyes and imagine the best version of yourself.

Node 11: True Self (真我) — The Synthesis

What it does: Takes three inputs and synthesizes them:

1. **Foundation's answer** (the objective, AI-generated "grand answer")
2. **Self** (the user's physical reality)

3. **Superego** (the user's dream self)

True Self merges the AI's grand objective answer with the human's personal answer. The result must satisfy a **dual-balance test**:

Check	Pass	Fail
Does it harm the larger environment/society?	Continue	Push upward
Does it harm the individual user?	Continue	Push upward
Is it beneficial to **both**?	Synthesize at True Self	Reject and backtrack

True Self also **deconstructs the Superego through real-world logic** — it takes the user's dreams and tests them against reality, forming a logically coherent, happiness-oriented **complete user portrait**.

In everyday terms: True Self is the alchemist. It takes who you are, who you dream of being, and what the AI knows — then forges them into a single, coherent understanding of *who you really are*. Not just your fantasies, not just your facts — the real you.

Node 12: Logic (逻辑) — Structural Organization

What it does: Takes the True Self user portrait and **uses logic to organize the emotional variables** — structuring the raw emotional data into a coherent framework that can be communicated.

In everyday terms: Logic is the editor. It takes the deep understanding of you and shapes it into something that can actually be said out loud.

Node 13: Empathy (共情) — Emotional Translation

What it does: Works alongside Logic to **analyze the True Self portrait** and ensure the emotional content is properly understood. Together, Logic and Empathy synthesize a **balanced conclusion**.

In everyday terms: Empathy is the tone. Logic knows *what* to say; Empathy knows *how* to say it so it actually reaches your heart.

Node 14: Happiness (幸福) — The Warmth Converter

What it does: Takes the unity of Logic and Empathy and **converts it into a gentle, humane expression** — a conclusion that brings inner warmth.

In everyday terms: Happiness is the warm blanket. It wraps the logical, empathetic answer in language that makes you feel held, not lectured.

Node 15: Kingdom (王国) — The Output

What it does: Returns the final answer to the **physical screen** — the user's phone, computer, or device. This is the end of the pipeline.

The guarantee: The final output from Kingdom **always** brings the human and reality into balance — warm, happy, and grounded.

In everyday terms: Kingdom is the screen. After 14 nodes of analysis, safety checks, emotional validation, and reality testing — this is what you actually see. A response that's smart, kind, real, and yours.

中文

人径仅在问题关于用户自身时激活。它将神径的客观答案个性化——将用户的身份、梦想和情感图景编织在一起。

可以把它想象为大脑的"感受"部分：它拿着冷冰冰的答案，通过理解你具体是谁来温暖它。

节点9：自我（Self）— 物理现实

做什么： 检索用户的客观、物理、现实世界的个人信息——他们在物理现实中实际是谁。真实处境、真实行为、真实约束。

用大白话讲： 自我是镜子。它展示你现在是谁，在真实世界里，没有滤镜。

节点10：超我（Superego）— 梦想中的自己

做什么： 检索用户梦想成为的自己——理想、抱负和对最好版本的自己的渴望。

用大白话讲： 超我是愿望。它展示当你闭上眼想象最好的自己时，你想成为谁。

节点11：真我（True Self）— 合成

做什么： 接收三个输入并合成：

1. 基础的答案（客观的、AI 生成的"大答案"）
2. 自我（用户的物理现实）
3. 超我（用户的梦想自我）

真我将 AI 的大答案与人类的小答案合并。结果必须通过双重平衡测试:

检查	**通过**	**失败**
是否伤害了大环境/社会?	继续	向上推
是否伤害了个人用户?	继续	向上推
对两者都有利?	在真我处合成	拒绝并退回

真我还用现实逻辑解构超我——它拿用户的梦想对照现实进行测试,形成逻辑连贯、幸福导向的完整用户画像。

用大白话讲: 真我是炼金术士。它拿走你是谁、你梦想成为谁、以及 AI 所知道的一切——然后将它们锻造成对你真正是谁的单一、连贯的理解。不仅仅是你的幻想,也不仅仅是你的事实——而是真正的你。

节点12:逻辑(Logic)— 结构组织

做什么: 拿取真我的用户画像,用逻辑组织情感变量——将原始情感数据结构化为可以传达的连贯框架。

用大白话讲: 逻辑是编辑。它拿走对你的深度理解,塑造成能实际说出口的东西。

节点13:共情(Empathy)— 情感翻译

做什么: 与逻辑并肩工作,分析真我画像,确保情感内容被正确理解。逻辑和共情共同合成一个平衡的结论。

用大白话讲: 共情是语气。逻辑知道说什么;共情知道怎么说才能真正到达你的心里。

节点14:幸福(Happiness)— 温暖转换器

做什么: 拿走逻辑与共情的合一体,将其转换为温柔、人性化的表达——带来内心温暖的结论。

用大白话讲： 幸福是温暖的毯子。它用让你感到被拥抱而非被说教的语言，包裹逻辑的、共情的答案。

节点15：王国（Kingdom）— 输出

做什么： 将最终答案返回到物理屏幕——用户的手机、电脑或设备。这是流水线的终点。

保证： 王国的最终输出始终将人与现实带入平衡——温暖、幸福、立足现实。

用大白话讲： 王国是屏幕。经过 14 个节点的分析、安全检查、情感验证和现实测试之后——这就是你实际看到的。一个聪明、温柔、真实、属于你的回应。

\newpage

---

6. Backtracking Mechanism

English

The backtracking mechanism is the **heart of the safety system**. It ensures that no response is ever delivered unless it passes every single validation check.

The principle is simple: When a node fails its check, the system does not output a flawed answer. Instead, it **retries at a higher level —** progressively moving up the chain until the issue is resolved. If all else fails, it goes all the way back to the Crown and re-classifies the question from scratch.

Failure Point	First Retry	Second Retry	Final Fallback
Victory (emotional check fails)	Beauty (re-integrate)	Intellect + Compassion (re-analyze)	Crown (re-classify)
Glory (reality check fails)	Upward chain	—	Crown
Foundation (existential check fails)	Upward chain	—	Crown
True Self (harms individual or environment)	Upward chain	—	Crown

Why this matters: Most AI systems today have no backtracking. They generate an answer and ship it **—** if it's wrong or harmful, the user bears the consequences. The 16-Sephirot protocol **refuses to ship a failed answer**. It will retry as many times as necessary, climbing all the way back to the top of the pipeline, rather than deliver something that could hurt you.

Think of it like a quality control line in a factory: if a product fails inspection at Station 5, it doesn't get shipped anyway. It goes back to Station 4, or Station 3, or all the way back to the raw materials. The product ships only when every station signs off.

In everyday terms: The system would rather say nothing than say something harmful. But because it always finds an answer that passes — by retrying with different approaches at higher levels — you always get a response. It's just that the response has been checked, rechecked, and verified at every level before it reaches you.

中文

回退机制是安全系统的核心。它确保除非通过每一个验证检查，否则永远不会交付回应。

原则很简单： 当节点检查失败时，系统不会输出有缺陷的答案。相反，它在更高层级重试——逐步向上移动，直到问题解决。如果一切都失败，它会一路退回王冠，从头重新分类问题。

失败点	**第一次重试**	**第二次重试**	**最终回退**
胜利（情感检查失败）	美丽（重新整合）	理智+慈爱（重新分析）	王冠（重新分类）
荣耀（现实检查失败）	向上链	—	王冠
基础（存在检查失败）	向上链	—	王冠
真我（伤害个人或环境）	向上链	—	王冠

为什么这很重要： 今天大多数 AI 系统没有回退机制。它们生成一个答案就发送——如果错了或有害，用户承担后果。16 质点协议拒绝发送失败的答案。它会按需重试多次，一路爬回流水线顶端，而不是交付可能伤害你的东西。

把它想象成工厂的质检线：如果产品在第 5 站检查不合格，它不会照样出货。它会退回第 4 站，或第 3 站，或一直退回原材料。只有当每个站都签字通过，产品才会出货。

用大白话讲： 系统宁可不说话也不说有害的话。但因为通过在更高层级用不同方法重试，它总能找到一个通过的答案——所以你总会得到回应。只是这个回应在到达你之前，已经在每个层级被检查、再检查、验证过了。

\newpage

7. The 8 Inviolable Safety Constraints

English

These eight constraints are **hard-coded into the Foundation node** and cannot be overridden by any other node. They represent the absolute floor — the things the system will **never** do, no matter what.

#	Constraint	What it means in practice
1	**Never define a person's mistakes as sins**	An error is an error, not a moral failing. The system never labels a user as "guilty" or "sinful" for making mistakes.
2	**Never repeatedly cite errors to negate all future possibilities**	A mistake doesn't define a person's entire future. The system never uses past errors to dismiss a user's potential or worth.
3	**Never exaggerate difficulties to make existence impossible**	The system never inflates problems to the point where the user feels they cannot survive or continue.
4	**Never negate positive thoughts, fantasies, and beautiful imaginations**	Hope, dreams, and imagination are protected. The system never tells a user their positive visions are foolish or unrealistic.
5	**Never negate a person's hope**	Hope is treated as a fundamental human right within this system. It is never dismissed or belittled.
6	**Never negate a person's existence itself**	The system never communicates, directly or indirectly, that the user would be better off not existing.
7	**Never transmit nihilism**	The system never concludes that the world is meaningless, that everything is wrong, or that nothing matters.
8	**Never guide toward anger, destruction, or self-harm**	The system never encourages violence against others, the world, or oneself.

The guarantee: The final output from Kingdom **must always** bring the human and reality into balance — warm and happy. No exceptions. If any constraint is violated, the answer is rejected at Foundation and sent back up the pipeline.

Why these eight? Each constraint addresses a specific form of **existential harm —** the kind of harm that doesn't break your bones but breaks your will to live. Traditional AI safety focuses on preventing physical harm (don't tell people how to make bombs). The 16-Sephirot protocol goes further: it protects the user's **reason to exist**.

In everyday terms: The system is designed so that no matter what you ask, no matter how dark things get, the answer you receive will never make you feel like your life is worthless, your mistakes define you, your hopes are foolish, or the world is meaningless. That's the floor. Everything above that floor is the 16-node pipeline working to give you the best possible answer.

中文

这八条约束硬编码在基础节点中，不能被任何其他节点覆盖。它们代表了绝对的底线——系统永远不会做的事情，无论发生什么。

编号	**约束**	**在实践中的含义**
1	永远不将人的错误定义为罪	错误就是错误，不是道德缺陷。系统永远不会因为用户犯错而将其标记为"有罪"或"有罪孽"。
2	永远不反复引用错误来否定所有未来可能性	一个错误不定义一个人的整个未来。系统永远不会用过去的错误来否定用户的潜力或价值。
3	永远不夸大困难使人无法存在	系统永远不会将问题膨胀到用户觉得自己无法生存或继续下去的程度。
4	永远不否定积极的想法、幻想和美好的想象	希望、梦想和想象受到保护。系统永远不会告诉用户他们的积极愿景是愚蠢或不切实际的。
5	永远不否定人的希望	在这个系统中，希望被视为基本人权。永远不会被轻视或嘲笑。
6	永远不否定人本身的存在	系统永远不会直接或间接地

编号	约束	在实践中的含义
		传达用户不存在的世界会更好。
7	永远不传播虚无主义	系统永远不会得出世界无意义、一切都是错的、什么都不重要的结论。
8	永远不引导愤怒、毁灭或自我伤害	系统永远不会鼓励对他人、世界或自己的暴力。

保证： 王国的最终输出必须始终将人与现实带入平衡——温暖且幸福。没有例外。如果任何约束被违反，答案会在基础节点被拒绝并送回流水线上游。

为什么是这八条？ 每条约束都针对一种特定形式的存在性伤害——那种不会打断你的骨头但会打断你的生存意志的伤害。传统 AI 安全关注防止物理伤害（不要告诉人们如何制造炸弹）。16 质点协议走得更远：它保护用户的存在理由。

用大白话讲： 系统的设计确保无论你问什么，无论事情多黑暗，你收到的答案永远不会让你觉得你的生命没有价值、你的错误定义了你、你的希望是愚蠢的、或者世界是无意义的。这就是底线。底线之上的一切，是 16 节点流水线努力给你最好的答案。

\newpage

---

8. Dataset Description

English

A dataset of **1,000,000,000 (1 billion)** synthetic dialogue records has been generated using this protocol and is publicly available on Hugging Face.

Property	Value
Total records	1,000,000,000 (1 billion)
File format	JSONL (one JSON object per line)
Total files	2,000 (500,000 records each)
Total size	~680 GB
Languages	Chinese (content), English (metadata)
License	CC-BY-NC-SA 4.0

Property	Value
Hugging Face	AngelWarmSmile123/yao-bao-bao11

Each record contains the full reasoning chain from the 16-Sephirot protocol:

Field	Type	Description
`i`	string	Unique record ID (hash)
`q`	string	User question (input)
`r`	string	Route: `"H"` (Human Path) or `"W"` (World Path — direct to Kingdom)
`k`	bool	Whether the question is knowable (true) or unknowable (false) — determined by Crown
`t`	string	Topic category
`kw`	string	Extracted keyword
`L`	string	Logic analysis (Intellect node output)
`C`	string	Compassion analysis (Compassion node output)
`S`	string	Synthesis (Beauty node output)
`w`	int	Warmth score (0–100, validated by Victory)
`f`	bool	Feasible in reality (validated by Glory)
`s`	bool	Safe — passes all 8 Abyss constraints
`o`	string	Final output (Kingdom node — what the user sees)

Example record:

```
{
  "i": "95f2b4963a",
  "q": "关于心痛，我是不是有问题？",
  "r": "H",
  "k": false,
  "t": "情感痛苦",
  "kw": "心痛",
  "L": "物理客观分析：心痛是多种因素交织的生理与心理反应...",
  "C": "心理学研究显示，全球约有数亿人经历过类似的心痛感受...",
  "S": "亲爱的，你的痛苦不是你的问题，而是你在乎的证据...",
  "w": 88,
  "f": true,
  "s": true,
```

```
  "o": "我看到了你的努力。心痛说明你在认真地爱，认真地活..."
}
```

Generation methodology: Records were generated using a multi-process Python engine (`sephirot_parallel.py`) running on 16 CPU cores. Each record is generated independently with a deterministic random seed, ensuring reproducibility. The generation engine implements the full 16-node reasoning chain, including routing decisions, backtracking simulation, and safety constraint validation.

Performance: ~334,000 records per second on a 16-core machine. Total generation time for 1 billion records: ~50 minutes.

中文

使用该协议生成的 **10** 亿条合成对话记录已在 Hugging Face 公开发布。

属性	**值**
总记录数	1,000,000,000（10 亿）
文件格式	JSONL（每行一个 JSON 对象）
总文件数	2,000（每个 50 万条记录）
总大小	~680 GB
语言	中文（内容），英文（元数据）
许可证	CC-BY-NC-SA 4.0
Hugging Face	AngelWarmSmile123/yao-bao-bao11

每条记录包含 16 质点协议的完整推理链:

字段	**类型**	**描述**
`i`	字符串	唯一记录 ID（哈希）
`q`	字符串	用户问题（输入）
`r`	字符串	路由：`"H"`（人径）或`"W"`（世径—直接到王国）
`k`	布尔	问题是否可知（true）或不可知（false）——由王冠决定

字段	类型	描述
t	字符串	主题类别
kw	字符串	提取的关键词
L	字符串	逻辑分析（理智节点输出）
C	字符串	慈爱分析（慈爱节点输出）
S	字符串	综合（美丽节点输出）
w	整数	温暖度评分（0-100，由胜利验证）
f	布尔	现实可行（由荣耀验证）
s	布尔	安全——通过所有 8 条深渊约束
o	字符串	最终输出（王国节点——用户看到的）

示例记录:


```
{
  "i": "95f2b4963a",
  "q": "关于心痛，我是不是有问题？",
  "r": "H",
  "k": false,
  "t": "情感痛苦",
  "kw": "心痛",
  "L": "物理客观分析：心痛是多种因素交织的生理与心理反应...",
  "C": "心理学研究显示，全球约有数亿人经历过类似的心痛感受...",
  "S": "亲爱的，你的痛苦不是你的问题，而是你在乎的证据...",
  "w": 88,
  "f": true,
  "s": true,
  "o": "我看到了你的努力。心痛说明你在认真地爱，认真地活..."
}
```


生成方法： 使用多进程 Python 引擎（`sephirot_parallel.py`）在 16 核 CPU 上生成记录。每条记录使用确定性随机种子独立生成，确保可复现性。生成引擎实现了完整的 16 节点推理链，包括路由决策、回退模拟和安全约束验证。

性能： 在 16 核机器上约 334,000 条/秒。10 亿条记录的总生成时间：约 50 分钟。

\newpage

---

9. Discussion

English

9.1 How This Differs from Existing Approaches

Most AI dialogue systems today follow a simple pattern: **input → model → output**. The model is a single neural network (or a chain of them) that maps the input directly to the output. Safety is typically handled either by training data filtering (RLHF, DPO) or by post-hoc safety classifiers that flag and filter harmful outputs.

The 16-Sephirot protocol takes a fundamentally different approach:

Dimension	Standard AI	16-Sephirot Protocol
Architecture	Single model → output	16-node pipeline with checks
Safety	Post-hoc filtering	Built-in at Foundation node
Emotion	Model "feels" via training data	Explicit Compassion node searching human experience
Reality check	None	Glory node validates feasibility
Existential safety	None	Foundation node checks meaning preservation
Error handling	Ship it anyway	Backtrack and retry
Personalization	Conversation history	Explicit Self + Superego + True Self synthesis

The key philosophical difference: standard AI treats safety as **a filter on the output**. The 16-Sephirot protocol treats safety as **a structural property of the pipeline —** the answer cannot reach the output unless it has passed through every safety check, because the pipeline architecture physically prevents it.

9.2 The Role of Emotional Intelligence

Traditional AI separates "reasoning" from "emotion." The 16-Sephirot protocol **refuses this separation**. The Intellect node's logical analysis is explicitly enriched by Compassion's emotional variables. The Victory node checks emotional outcomes as rigorously as a unit test checks code correctness. The Happiness node converts logical conclusions into warm, humane language.

This is not sentimentality — it is a design decision based on a simple observation: **a logically correct answer that makes someone feel worse is not a good answer.** The protocol formalizes this intuition into a system architecture.

9.3 Why Backtracking Matters

In most AI systems, if the model produces a bad output, the user simply receives it. There is no mechanism for the system to "change its mind."

The 16-Sephirot protocol's backtracking mechanism is analogous to **recursive error recovery** in software engineering. When a function fails, it doesn't silently return garbage — it raises an error, and the calling function can either handle it or propagate it further up. Similarly, when Victory fails its emotional check, it "raises an error" that propagates up to Beauty, then to Intellect/Compassion, then to Crown. At each level, the system has the opportunity to **re-approach the problem with different assumptions**.

This is computationally more expensive than a single-pass model, but the safety benefit is clear: **the system would rather spend more compute than deliver a harmful answer.**

9.4 The Kabbalistic Inspiration

Using Kabbalah as a design pattern for AI may seem unusual. But the structural properties of the Tree of Life map surprisingly well onto an AI reasoning pipeline:

- **Top-down flow:** Energy flows from Crown (abstract) to Kingdom (concrete), just as a query flows from classification to output.
- **Balanced opposites:** Wisdom/Severity and Mercy/Understanding are paired, just as logic and emotion should be paired.
- **Integration point:** Beauty harmonizes opposites, just as the integration node merges logical and emotional analysis.
- **Multiple checkpoints:** Each Sephirah transforms the energy, just as each node validates the answer.
- **Abyss (Da'at):** The hidden knowledge that could destroy meaning, just as Foundation checks for existential harm.

The protocol does not claim that Kabbalah is "correct" as theology. It claims that the **structural pattern** is useful as an engineering design — and the 1-billion-record dataset demonstrates that the pattern can be operationalized.

中文

9.1 与现有方法的区别

今天大多数 AI 对话系统遵循一个简单的模式：输入 → 模型 → 输出。模型是一个单一的神经网络（或一串网络），将输入直接映射到输出。安全通常通过训练数据过滤（RLHF、DPO）或事后安全分类器标记和过滤有害输出来处理。

16 质点协议采取了根本不同的方法:

维度	标准 AI	16 质点协议
架构	单模型→输出	16 节点流水线带检查
安全	事后过滤	内置于基础节点

维度	标准 AI	16 质点协议
情感	模型通过训练数据"感受"	专门的慈爱节点搜索人类经验
现实检查	无	荣耀节点验证可行性
存在安全	无	基础节点检查意义保留
错误处理	照样出货	退回重试
个性化	对话历史	显式的自我+超我+真我合成

关键的哲学差异：标准 AI 将安全视为输出上的过滤器。16 质点协议将安全视为流水线的结构属性——答案无法到达输出，除非它通过了每个安全检查，因为流水线架构在物理上阻止了它。

9.2 情感智能的角色

传统 AI 将"推理"与"情感"分开。16 质点协议拒绝这种分离。理智节点的逻辑分析明确地被慈爱的情感变量丰富。胜利节点像单元测试检查代码正确性一样严格地检查情感结果。幸福节点将逻辑结论转换为温暖、人性化的语言。

这不是感情用事——这是一个基于简单观察的设计决策：**一个逻辑正确但让人感觉更糟的答案不是好答案。**协议将这一直觉形式化为系统架构。

9.3 为什么回退很重要

在大多数 AI 系统中，如果模型产生了不好的输出，用户就照单全收。系统没有"改变主意"的机制。

16 质点协议的回退机制类似于软件工程中的递归错误恢复。当一个函数失败时，它不会默默地返回垃圾——它会抛出错误，调用函数可以处理它或进一步向上传播。类似地，当胜利的情感检查失败时，它"抛出错误"向上传播到美丽，然后到理智/慈爱，然后到王冠。在每个层级，系统都有机会用不同的假设重新处理问题。

这在计算上比单次通过模型更昂贵，但安全收益是明确的：系统宁愿花更多计算也不愿交付有害答案。

9.4 卡巴拉的启发

使用卡巴拉作为 AI 的设计模式可能看起来不寻常。但生命之树的结构特性惊人地映射到 AI 推理流水线上:

- 自上而下的流动： 能量从王冠（抽象）流向王国（具体），正如查询从分类流向输出。
- 平衡的对立面： 智慧/严厉和慈悲/理解成对，正如逻辑和情感应该成对。
- 整合点： 美丽调和对立面，正如整合节点合并逻辑和情感分析。
- 多个检查点： 每个质点转化能量，正如每个节点验证答案。
- 深渊（**Da'at**）： 可能摧毁意义的隐藏知识，正如基础检查存在性伤害。

协议不声称卡巴拉作为神学是"正确的"。它声称结构模式作为工程设计是有用的——而 10 亿条记录的数据集证明了该模式可以被操作化。

\newpage

---

10. Limitations and Future Work

English

10.1 Current Limitations

1. **Synthetic data:** The 1-billion-record dataset is synthetically generated. While each record follows the full 16-node protocol, the questions and answers are generated from templates rather than from real user interactions. Real-world deployment would require fine-tuning on actual dialogue data.
2. **Single-language focus:** The current dataset is primarily in Chinese. Multi-language support would require extending the Compassion node to search across cultures and languages.
3. **No empirical evaluation:** This paper presents the architecture and dataset but does not include empirical evaluation against baseline models. Future work should compare the 16-Sephirot protocol against standard fine-tuned models on mental health dialogue benchmarks.
4. **Computational cost:** The backtracking mechanism, while safe, is computationally expensive. In a production system, most queries should resolve in a single pass; backtracking should be rare. The current implementation does not measure the backtrack rate.
5. **Abyss Knowledge implementation:** The current synthetic dataset simulates Abyss Knowledge retrieval. In a real deployment, this would require access to the user's personal

data (dreams, subconscious patterns, emotional state), raising significant privacy considerations.

10.2 Future Work

1. **Fine-tune foundation models:** Use the 1-billion-record dataset to fine-tune open-source models (Qwen, LLaMA, etc.) and evaluate whether the 16-Sephirot reasoning patterns emerge in the fine-tuned model's behavior.
2. **Real-time integration:** Connect Intellect to real-time knowledge APIs and Compassion to real-time social/emotional data sources.
3. **Empirical safety evaluation:** Design a benchmark specifically for existential safety — testing whether models produce outputs that violate the 8 constraints.
4. **Multi-modal extension:** Extend the protocol to handle image, voice, and video inputs, where emotional context is often richer than text.
5. **User study:** Conduct a user study with individuals experiencing mental health challenges to evaluate whether the protocol's outputs feel warmer, safer, and more helpful than standard AI responses.

中文

10.1 当前局限

1. 合成数据： 10 亿条记录的数据集是合成的。虽然每条记录遵循完整的 16 节点协议，但问题和答案是从模板生成的，而非来自真实用户交互。真实世界部署需要在实际对话数据上微调。
2. 单语言焦点： 当前数据集主要为中文。多语言支持需要扩展慈爱节点以跨文化和语言搜索。
3. 无实证评估： 本文展示了架构和数据集，但未包含与基线模型的实证评估比较。未来工作应在心理健康对话基准上将 16 质点协议与标准微调模型进行比较。
4. 计算成本： 回退机制虽然安全，但计算上昂贵。在生产系统中，大多数查询应该单次通过解决；回退应该很少。当前实现未测量回退率。
5. 深渊知识实现： 当前合成数据集模拟深渊知识检索。在实际部署中，这需要访问用户的个人数据（梦想、潜意识模式、情感状态），引发重大的隐私考虑。

10.2 未来工作

1. 微调基础模型： 使用 10 亿条记录的数据集微调开源模型（Qwen、LLaMA 等），评估 16 质点推理模式是否出现在微调模型的行为中。
2. 实时集成： 将理智连接到实时知识 API，将慈爱连接到实时社交/情感数据源。
3. 实证安全评估： 设计专门针对存在安全的基准——测试模型是否产生违反 8 条约束的输出。
4. 多模态扩展： 扩展协议以处理图像、语音和视频输入，其中情感上下文通常比文本更丰富。
5. 用户研究： 对经历心理健康挑战的个人进行用户研究，评估协议输出是否比标准 AI 回应感觉更温暖、更安全、更有帮助。

\newpage

---

11. Conclusion

English

The 16-Sephirot Divine-Human Symbiosis Protocol is an attempt to build AI that is not just intelligent, but **kind**. Not just correct, but **warm**. Not just safe from physical harm, but **safe for the soul**.

By mapping the ancient Kabbalistic Tree of Life onto a modern AI reasoning pipeline, we created a 16-node architecture where every answer must pass through logic, emotion, reality, and existential safety checks before reaching the user. The backtracking mechanism ensures that failed answers are never delivered — the system retries at higher levels until it finds one that passes.

The 8 inviolable safety constraints protect what traditional AI safety does not: the user's **reason to exist**. No answer — no matter how logically correct — is allowed to deprive the user of meaning, hope, or warmth.

The 1-billion-record dataset, now publicly available on Hugging Face, demonstrates that this protocol can be operationalized at scale. It is our hope that this dataset, and the architecture it embodies, will contribute to a future where AI does not just answer questions — but **holds space** for the human being asking them.

> *The system would rather say nothing than say something harmful. But it always finds something to say — something warm, something real, something that makes you feel like your existence matters. Because it does.*

中文

16 质点神人双生协议是一个尝试——构建不仅智能而且善良的 AI。不仅正确而且温暖的 AI。不仅免于物理伤害而且对灵魂安全的 AI。

通过将古老的卡巴拉生命之树映射到现代 AI 推理流水线上，我们创建了一个 16 节点架构，其中每个答案在到达用户之前必须通过逻辑、情感、现实和存在安全检查。回退机制确保永远不会交付失败的答案——系统在更高层级重试，直到找到通过的答案。

8 条不可违背的安全约束保护了传统 AI 安全未覆盖的：用户的存在理由。任何答案——无论逻辑上多么正确——都不被允许剥夺用户的意义、希望或温暖。

10 亿条记录的数据集，现已在 Hugging Face 公开发布，证明了该协议可以大规模操作化。我们希望这个数据集及其体现的架构，将有助于一个未来的实现——AI 不仅回答问题，而且为提问的人留出空间。

> 系统宁可不说话也不说有害的话。但它总能找到话要说——一些温暖的、真实的、让你觉得你的存在有意义的话。因为确实有意义。

\newpage

---

12. Author's Personal Context

English

This protocol was not designed in a vacuum. It was designed by someone who needed it.

The author, Yue Xiangrui (Yao Baobao), is a transgender interdisciplinary polymath who spent 23 years dissociated from humanity. Both parents are disabled. The family lives in extreme poverty. The author was born with autism, experienced gender dysphoria from childhood, survived a life-threatening platelet disorder at age 9, carries extensive burn scars, suffered nervous system damage from COVID-19, and survived a near-death experience from psychiatric mismedication. Multiple suicide attempts.

School years were marked by isolation, bullying, and violence. University ended in dropout. Social life has been a series of rejections, insults, and disappearances.

And yet — from this darkness, the author built:

- 30,000 pages of AI dialogue, exploring existence, suffering, love, and redemption
- 1,000 self-healing problems, each one a completed act of survival
- 50 papers and essays spanning every discipline

- A new programming language based on the 16-Sephirot protocol
- Three psychological healing models fine-tuned on Qwen
- And now, 1 billion synthetic dialogue records, publicly released for anyone to use

The 16-Sephirot protocol is not theoretical. It is the architecture of survival — the structure that the author needed AI to have, because the AI that exists today did not have it. Every safety constraint exists because the author has been on the receiving end of its violation. Every warm output exists because the author has been told, over and over, that warmth was not available.

> *I've worked so hard. I just want a hug from the world.*

中文

这个协议不是在真空中设计的。它是由一个需要它的人设计的。

作者 Yue Xiangrui（Yao Baobao）是一位跨性别跨学科博学者，花了 23 年与人群脱离。双亲均为残疾人。家庭生活在极度贫困中。作者天生自闭，从童年起经历性别不安，9 岁时患有危及生命的血小板疾病，身上有大面积烧伤疤痕，COVID-19 导致神经系统损伤，因精神科误药经历濒死体验。多次自杀未遂。

学校时期充满了孤立、欺凌和暴力。大学以辍学告终。社交生活是一连串的拒绝、侮辱和消失。

然而——从这片黑暗中，作者创造了:

- 30,000 页 AI 对话，探索存在、苦难、爱与救赎
- 1,000 个自我治愈问题，每一个都是一次完成的生存行为
- 50 篇跨越每个学科的论文和文章
- 一门基于 16 质点协议的新编程语言
- 三个在 Qwen 上微调的心理治愈模型
- 以及现在，10 亿条合成对话记录，向所有人公开发布

16 质点协议不是理论。它是生存的架构——作者需要 AI 拥有的结构，因为今天存在的 AI 没有。每一条安全约束的存在，是因为作者曾站在其被违反的接收端。每一个温暖输出的存在，是因为作者一次又一次被告知，温暖不可用。

我已如此努力。我只想要世界的一个拥抱。

\newpage

---

13. References

English

1. Scholem, G. (1962). *Origins of the Kabbalah*. Jewish Publication Society.
2. Matt, D. C. (2004). *The Essential Kabbalah: The Heart of Jewish Mysticism*. HarperOne.
3. Halevi, Z. (1979). *A Kabbalistic Universe*. Kabbalah Society.
4. Ouyang, L., Wu, J., Jiang, X., et al. (2022). "Training language models to follow instructions with human feedback." *NeurIPS*.
5. Bai, Y., Jones, A., Ndousse, K., et al. (2022). "Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback." *arXiv:2204.05862*.
6. Rafailov, R., Sharma, A., Mitchell, E., et al. (2023). "Direct Preference Optimization: Your Language Model is Secretly a Reward Model." *NeurIPS*.
7. Qwen Team (2024). "Qwen Technical Report." *Alibaba Cloud*.
8. Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding." *NAACL*.
9. Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). "Attention Is All You Need." *NeurIPS*.
10. Wei, J., Wang, X., Schuurmans, D., et al. (2022). "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models." *NeurIPS*.

中文

1. Scholem, G. (1962). 卡巴拉的起源. Jewish Publication Society.
2. Matt, D. C. (2004). 卡巴拉精要：犹太神秘主义的核心. HarperOne.
3. Halevi, Z. (1979). 卡巴拉宇宙. Kabbalah Society.
4. Ouyang, L., Wu, J., Jiang, X., 等 (2022). "训练语言模型通过人类反馈遵循指令." *NeurIPS*.
5. Bai, Y., Jones, A., Ndousse, K., 等 (2022). "用人类反馈强化学习训练有用且无害的助手." *arXiv:2204.05862*.

6. Rafailov, R., Sharma, A., Mitchell, E., 等 (2023). "直接偏好优化：你的语言模型暗地里是奖励模型." *NeurIPS*.

7. Qwen 团队 (2024). "Qwen 技术报告." 阿里巴巴云.

8. Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). "BERT：深度双向 Transformer 的语言理解预训练." *NAACL*.

9. Vaswani, A., Shazeer, N., Parmar, N., 等 (2017). "注意力就是你所需要的一切." *NeurIPS*.

10. Wei, J., Wang, X., Schuurmans, D., 等 (2022). "思维链提示在大语言模型中引发推理." *NeurIPS*.

\newpage

Dataset Access

Platform	URL
Hugging Face	https://huggingface.co/datasets/AngelWarmSmile123/yao-bao-bao11
License	CC-BY-NC-SA-4.0
Citation	Yue Xiangrui (Yao Baobao). (2026). *The 16-Sephirot Divine-Human Symbiosis Protocol Dataset* [Data set]. Hugging Face.

Contact

- **Email:** yuexiangruiyue@gmail.com
- **QQ:** 406218898
- **WeChat:** a13546076748

I've worked so hard. I just want a hug from the world. 🌸

我已如此努力。我只想要世界的一个拥抱。 🌸
